

## ▼ Ungraded Lab: Model Analysis with TFX Evaluator

Now that you've used TFMA as a standalone library in the previous lab, you will now see how it is used by TFX with its [Evaluator](#) component. This component comes after your `Trainer` run and it checks if your trained model meets the minimum required metrics and also compares it with previously generated models.

You will go through a TFX pipeline that prepares and trains the same model architecture you used in the previous lab. As a reminder, this is a binary classifier to be trained on the [Census Income dataset](#). Since you're already familiar with the earlier TFX components, we will just go over them quickly but we've placed notes on where you can modify code if you want to practice or produce a better result.

Let's begin!

*Credits: Some of the code and discussions are based on the TensorFlow team's [official tutorial](#).*

### ▼ Setup

#### ▼ Install TFX

```
!pip install -U tfx==0.30
```

*Note: In Google Colab, you need to restart the runtime at this point to finalize updating the packages you just installed. You can do so by clicking the `Restart Runtime` at the end of the output cell above (after installation), or by selecting `Runtime > Restart Runtime` in the Menu bar. **Please do not proceed to the next section without restarting.** You can also ignore the errors about version incompatibility of some of the bundled packages because we won't be using those in this notebook.*

#### ▼ Imports

```
import os
import pprint

from tfx.components import CsvExampleGen
from tfx.components import ExampleValidator
from tfx.components import SchemaGen
from tfx.components import StatisticsGen
from tfx.components import Transform
from tfx.components import Tuner
```

```

from tfx.components import Trainer
from tfx.components import Evaluator

import tensorflow as tf
import tensorflow_model_analysis as tfma

from tfx.orchestration.experimental.interactive.interactive_context import InteractiveContext

tf.get_logger().propagate = False
tf.get_logger().setLevel('ERROR')
pp = pprint.PrettyPrinter()

```

## ▼ Set up pipeline paths

```

# Location of the pipeline metadata store
_pipeline_root = './pipeline/'

# Directory of the raw data files
_data_root = './data/census'

_data_filepath = os.path.join(_data_root, "data.csv")

# Create the TFX pipeline files directory
!mkdir {_pipeline_root}

# Create the dataset directory
!mkdir -p {_data_root}

```

## ▼ Download and prepare the dataset

Here, you will download the training split of the [Census Income Dataset](#). This is twice as large as the test dataset you used in the previous lab.

```

# Define filename and URL
TAR_NAME = 'lab_2_data.tar.gz'
DATA_PATH = f'https://github.com/https-deeplearning-ai/MLLP-public/raw/main/course3/week4-ungraded-lab2\_data.tar.gz'

# Download dataset
!wget -nc {DATA_PATH}

# Extract archive
!tar xvzf {TAR_NAME}

# Delete archive
!rm {TAR_NAME}

```

```

--2021-08-02 11:22:47-- https://github.com/https-deeplearning-ai/MLLP-public/raw/main/course3/week4-ungraded-lab2\_data.tar.gz
Resolving github.com (github.com)... 52.69.186.44
Connecting to github.com (github.com)|52.69.186.44|:443... connected.
HTTP request sent, awaiting response... 302 Found

```

```
Location: https://raw.githubusercontent.com/https-deeplearning-ai/MLLP-public/main/course3/w
--2021-08-02 11:22:47-- https://raw.githubusercontent.com/https-deeplearning-ai/MLLP-public
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.110.133, 185.199.
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.110.133|:443...
HTTP request sent, awaiting response... 200 OK
Length: 418898 (409K) [application/octet-stream]
Saving to: 'lab_2_data.tar.gz'
```

```
lab_2_data.tar.gz 100%[=====>] 409.08K --.-KB/s in 0.04s
```

```
2021-08-02 11:22:48 (9.27 MB/s) - 'lab_2_data.tar.gz' saved [418898/418898]
```

```
./data/census/data.csv
```



Take a quick look at the first few rows.

```
# Preview dataset
```

```
!head {_data_filepath}
```

```
age,workclass,fnlwgt,education,education-num,marital-status,occupation,relationship,race,sex
39,State-gov,77516,Bachelors,13,Never-married,Adm-clerical,Not-in-family,White,Male,2174,0,4
50,Self-emp-not-inc,83311,Bachelors,13,Married-civ-spouse,Exec-managerial,Husband,White,Male
38,Private,215646,HS-grad,9,Divorced,Handlers-cleaners,Not-in-family,White,Male,0,0,40,Unit
53,Private,234721,11th,7,Married-civ-spouse,Handlers-cleaners,Husband,Black,Male,0,0,40,Unit
28,Private,338409,Bachelors,13,Married-civ-spouse,Prof-specialty,Wife,Black,Female,0,0,40,Cul
37,Private,284582,Masters,14,Married-civ-spouse,Exec-managerial,Wife,White,Female,0,0,40,Uni
49,Private,160187,9th,5,Married-spouse-absent,Other-service,Not-in-family,Black,Female,0,0,1
52,Self-emp-not-inc,209642,HS-grad,9,Married-civ-spouse,Exec-managerial,Husband,White,Male,0
31,Private,45781,Masters,14,Never-married,Prof-specialty,Not-in-family,White,Female,14084,0,!
```



## ▼ TFX Pipeline

### ▼ Create the InteractiveContext

As usual, you will initialize the pipeline and use a local SQLite file for the metadata store.

```
# Initialize InteractiveContext
```

```
context = InteractiveContext(pipeline_root=pipeline_root)
```

```
WARNING:absl:InteractiveContext metadata_connection_config not provided: using SQLite ML Met
```



### ▼ ExampleGen

You will start by ingesting the data through `CsvExampleGen`. The code below uses the default 2:1 train-eval split (i.e. 33% of the data goes to eval) but feel free to modify if you want. You can review splitting techniques [here](#).

```
# Run CsvExampleGen
example_gen = CsvExampleGen(input_base=_data_root)
context.run(example_gen)
```

WARNING:apache\_beam.runners.interactive.interactive\_environment:Dependencies required for In  
 WARNING:root:Make sure that locally built Python SDK docker image has Python 3.7 interpreter  
 WARNING:apache\_beam.io.tfrecordio:Couldn't find python-snappy so the implementation of \_TFRe

▼ **ExecutionResult** at 0x7fb4587a6a10

```
.execution_id      1
.component         ► CsvExampleGen at 0x7fb3e561bb90
.component.inputs   {}
.component.outputs  ['examples'] ► Channel of type 'Examples' (1 artifact) at
                                0x7fb3e561bf50
```

```
# Print split names and URI
artifact = example_gen.outputs['examples'].get()[0]
print(artifact.split_names, artifact.uri)
```

```
["train", "eval"] ./pipeline/CsvExampleGen/examples/1
```

## ▼ StatisticsGen

You will then compute the statistics so it can be used by the next components.

```
# Run StatisticsGen
statistics_gen = StatisticsGen(
    examples=example_gen.outputs['examples'])
context.run(statistics_gen)
```

WARNING:root:Make sure that locally built Python SDK docker image has Python 3.7 interpreter

▼ **ExecutionResult** at 0x7fb3e4ce9310

```
.execution_id      2
.component         ► StatisticsGen at 0x7fb3e45c0610
.component.inputs   ['examples'] ► Channel of type 'Examples' (1 artifact) at
                                0x7fb3e561bf50
.component.outputs  ['statistics'] ► Channel of type 'ExampleStatistics' (1 artifact) at
                                0x7fb3e45c0610
```

You can look at the visualizations below if you want to explore the data some more.

```
# Visualize statistics
context.show(statistics_gen.outputs['statistics'])
```

Categorical Features (8)						Chart to show
count	missing	unique	top	freq top	avg str len	Standard
						<input type="checkbox"/> log <input type="checkbox"/> expand



## ▼ SchemaGen

You can then infer the dataset schema with [SchemaGen](#). This will be used to validate incoming data to ensure that it is formatted correctly.

```
# Run SchemaGen
schema_gen = SchemaGen(
    statistics=statistics_gen.outputs['statistics'])
context.run(schema_gen)
```

```
▼ ExecutionResult at 0x7fb3e4919210
  .execution_id      3
  .component         ► SchemaGen at 0x7fb3e3e1c150
  .component.inputs  ['statistics'] ► Channel of type 'ExampleStatistics' (1 artifact) at
                                0x7fb3e45c0190
  .component.outputs ['schema'] ► Channel of type 'Schema' (1 artifact) at
                                0x7fb3e3410000
```

For simplicity, you will just accept the inferred schema but feel free to modify with the [TFDV API](#) if you want.

```
# Visualize the inferred Schema
context.show(schema_gen.outputs['schema'])
```

	Type	Presence	Valency	Domain
Feature name				
'education'	STRING	required		'education'
'marital-status'	STRING	required		'marital-status'
'native-country'	STRING	required		'native-country'
'occupation'	STRING	required		'occupation'
'race'	STRING	required		'race'
'relationship'	STRING	required		'relationship'
'sex'	STRING	required		'sex'
'workclass'	STRING	required		'workclass'
'age'	INT	required		-
'capital-gain'	INT	required		-
'capital-loss'	INT	required		-
'education-num'	INT	required		-
'fnlwgt'	INT	required		-
'hours-per-week'	INT	required		-
'label'	INT	required		-

## Values

Domain	
<b>'education'</b>	'10th', '11th', '12th', '1st-4th', '5th-6th', '7th-8th', '9th', 'Assoc-acdm', 'Assoc-voc', 'Bachelors', 'Doctorate', 'HS-grad', 'Masters', 'Preschool', 'Prof-school', 'Some-college'
<b>'marital-</b>	'Divorced', 'Married-AF-spouse', 'Married-civ-spouse', 'Married-spouse-absent',

Next, run `ExampleValidator` to check if there are anomalies in the data.

```
# Run ExampleValidator
```



```
▼ ExecutionResult at 0x7fb3e3e5a5d0
  .execution_id      4
  .component         ► ExampleValidator at 0x7fb3e3f10b90
  .component.inputs  ['statistics'] ► Channel of type 'ExampleStatistics' (1 artifact) at
                                0x7fb3e45c0190
```

If you just used the inferred schema, then there should not be any anomalies detected. If you modified the schema, then there might be some results here and you can again use TFDV to modify or relax constraints as needed.

In actual deployments, this component will also help you understand how your data evolves over time and identify data errors. For example, the first batches of data that you get from your users might conform to the schema but it might not be the case after several months. This component will detect that and let you know that your model might need to be updated.

```
# Check results
context.show(example_validator.outputs['anomalies'])
```

**Artifact at ./pipeline/ExampleValidator/anomalies/4**

**'train' split:**

```
/usr/local/lib/python3.7/dist-packages/tensorflow_data_validation/utils/display_util.py:217:
pd.set_option('max_colwidth', -1)
```

**No anomalies found.**

**'eval' split:**

**No anomalies found.**

## ▼ Transform

Now you will perform feature engineering on the training data. As shown when you previewed the CSV earlier, the data is still in raw format and cannot be consumed by the model just yet. The transform code in the following cells will take care of scaling your numeric features and one-hot encoding your categorical variables.

*Note: If you're running this exercise for the first time, we advise that you leave the transformation code as is. After you've gone through the entire notebook, then you can modify these for practice and see what results you get. Just make sure that your model builder code in the Trainer component will also reflect those changes if needed. For example, removing a feature here should also remove an input layer for that feature in the model.*

```
# Set the constants module filename
_census_constants_module_file = 'census_constants.py'
```

```

%%writefile {_census_constants_module_file}

# Features with string data types that will be converted to indices
VOCAB_FEATURE_DICT = {
    'education': 16, 'marital-status': 7, 'occupation': 15, 'race': 5,
    'relationship': 6, 'workclass': 9, 'sex': 2, 'native-country': 42
}

# Numerical features that are marked as continuous
NUMERIC_FEATURE_KEYS = ['fnlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']

# Feature that can be grouped into buckets
BUCKET_FEATURE_DICT = {'age': 4}

# Number of out-of-vocabulary buckets
NUM_OOV_BUCKETS = 5

# Feature that the model will predict
LABEL_KEY = 'label'

# Utility function for renaming the feature
def transformed_name(key):
    return key + '_xf'

    Writing census_constants.py

# Set the transform module filename
_census_transform_module_file = 'census_transform.py'

%%writefile {_census_transform_module_file}

import tensorflow as tf
import tensorflow_transform as tft

# import constants from cells above
import census_constants

# Unpack the contents of the constants module
_NUMERIC_FEATURE_KEYS = census_constants.NUMERIC_FEATURE_KEYS
_VOCAB_FEATURE_DICT = census_constants.VOCAB_FEATURE_DICT
_BUCKET_FEATURE_DICT = census_constants.BUCKET_FEATURE_DICT
_NUM_OOV_BUCKETS = census_constants.NUM_OOV_BUCKETS
_LABEL_KEY = census_constants.LABEL_KEY
_transformed_name = census_constants.transformed_name

# Define the transformations
def preprocessing_fn(inputs):
    """tf.transform's callback function for preprocessing inputs.
    Args:
        inputs: map from feature keys to raw not-yet-transformed features.
    Returns:
        Map from string feature key to transformed feature operations.
    """

```

```

# Initialize outputs dictionary
outputs = {}

# Scale these features to the range [0,1]
for key in _NUMERIC_FEATURE_KEYS:
    scaled = tft.scale_to_0_1(inputs[key])
    outputs[_transformed_name(key)] = tf.reshape(scaled, [-1])

# Convert strings to indices and convert to one-hot vectors
for key, vocab_size in _VOCAB_FEATURE_DICT.items():
    indices = tft.compute_and_apply_vocabulary(inputs[key], num_oov_buckets=_NUM_OOV_BUCKETS)
    one_hot = tf.one_hot(indices, vocab_size + _NUM_OOV_BUCKETS)
    outputs[_transformed_name(key)] = tf.reshape(one_hot, [-1, vocab_size + _NUM_OOV_BUCKETS])

# Bucketize this feature and convert to one-hot vectors
for key, num_buckets in _BUCKET_FEATURE_DICT.items():
    indices = tft.bucketize(inputs[key], num_buckets)
    one_hot = tf.one_hot(indices, num_buckets)
    outputs[_transformed_name(key)] = tf.reshape(one_hot, [-1, num_buckets])

# Cast label to float
outputs[_transformed_name(_LABEL_KEY)] = tf.cast(inputs[_LABEL_KEY], tf.float32)

return outputs

Writing census_transform.py

```

Now, we pass in this feature engineering code to the **Transform** component and run it to transform your data.

```

# Run the Transform component
transform = Transform(
    examples=example_gen.outputs['examples'],
    schema=schema_gen.outputs['schema'],
    module_file=os.path.abspath(_census_transform_module_file))
context.run(transform, enable_cache=False)

```

```

ERROR:absl:udf_utils.get_fn {'module_file': None, 'module_path': 'census_transform@./pipeline
WARNING:root:This output type hint will be ignored and not used for type-checking purposes.
WARNING:root:This output type hint will be ignored and not used for type-checking purposes.
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring return_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring return_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring return_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring return_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring return_type hint: <class 'NoneType'>
WARNING:apache_beam.typehints.typehints:Ignoring send_type hint: <class 'NoneType'>

```

You can see a sample result for one row with the code below. Notice that the numeric features are indeed scaled and the categorical features are now one-hot encoded.

```

| .execution id          5 |
# Get the URI of the output artifact representing the transformed examples
train_uri = os.path.join(transform.outputs['transformed_examples'].get()[0].uri, 'Split-train')

# Get the list of files in this directory (all compressed TFRecord files)
tfrecord_filenames = [os.path.join(train_uri, name)
                       for name in os.listdir(train_uri)]

# Create a `TFRecordDataset` to read these files
dataset = tf.data.TFRecordDataset(tfrecord_filenames, compression_type="GZIP")

# Decode the first record and print output
for tfrecord in dataset.take(1):
    serialized_example = tfrecord.numpy()
    example = tf.train.Example()
    example.ParseFromString(serialized_example)
    pp.pprint(example)

        value: 0.0
        value: 0.0
        value: 0.0
      }
    }
  }
  feature {
    key: "relationship_xf"
    value {
      float_list {
        value: 0.0
        value: 1.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
      }
    }
  }
}

```

```

    },
  },
  feature {
    key: "sex_xf"
    value {
      float_list {
        value: 1.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
      }
    }
  }
  feature {
    key: "workclass_xf"
    value {
      float_list {
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 1.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
        value: 0.0
      }
    }
  }
}

```

As you already know, the `Transform` component not only outputs the transformed examples but also the transformation graph. This should be used on all inputs when your model is deployed to ensure that it is transformed the same way as your training data. Else, it can produce training-serving skew which leads to noisy predictions.

The `Transform` component stores related files in its `transform_graph` output and it would be good to quickly review its contents before we move on to the next component. As shown below, the URI of this output points to a directory containing three subdirectories.

```

# Get URI and list subdirectories
graph_uri = transform.outputs['transform_graph'].get()[0].uri
os.listdir(graph_uri)

['transform_fn', 'metadata', 'transformed_metadata']

```

- The `transformed_metadata` subdirectory contains the schema of the preprocessed data.
- The `transform_fn` subdirectory contains the actual preprocessing graph.

- The `metadata` subdirectory contains the schema of the original data.

## ▼ Trainer

Next, you will now build the model to make your predictions. As mentioned earlier, this is a binary classifier where the label is 1 if a person earns more than 50k USD and 0 if less than or equal. The model used here uses the [wide and deep model](#) as reference but feel free to modify after you've completed the exercise. Also for simplicity, the hyperparameters (e.g. number of hidden units) have been hardcoded but feel free to use a `Tuner` component as you did in Week 1 if you want to get some practice.

As a reminder, it is best to start from `run_fn()` when you're reviewing the module file below. The `Trainer` component looks for that function first. All other functions defined in the module are just utility functions for `run_fn()`.

Another thing you will notice below is the `_get_serve_tf_examples_fn()` function. This is tied to the `serving_default` [signature](#) which makes it possible for you to just pass in raw features when the model is served for inference. You have seen this in action in the previous lab. This is done by decorating the enclosing function, `serve_tf_examples_fn()`, with [tf.function](#). This indicates that the computation will be done by first tracing a [Tensorflow graph](#). You will notice that this function uses `model.tft_layer` which comes from `transform_graph` output. Now when you call the `.get_concrete_function()` method on this `tf.function` in `run_fn()`, you are creating the graph that will be used in later computations. This graph is used whenever you pass in an `examples` argument pointing to a Tensor with `tf.string` dtype. That matches the format of the serialized batches of data you used in the previous lab.

```
# Declare trainer module file
_census_trainer_module_file = 'census_trainer.py'

%%writefile {_census_trainer_module_file}

from typing import List, Text

import tensorflow as tf
import tensorflow_transform as tft
from tensorflow_transform.tf_metadata import schema_utils

from tfx.components.trainer.fn_args_utils import DataAccessor, FnArgs
from tfx_bsl.public.tfxio import TensorFlowDatasetOptions

# import same constants from transform module
import census_constants

# Unpack the contents of the constants module
_NUMERIC_FEATURE_KEYS = census_constants.NUMERIC_FEATURE_KEYS
_VOCAB_FEATURE_DICT = census_constants.VOCAB_FEATURE_DICT
_BUCKET_FEATURE_DICT = census_constants.BUCKET_FEATURE_DICT
_NUM_CLASS_BUCKETS = census_constants.NUM_CLASS_BUCKETS
```

```

_NUM_OUV_BUCKETS = census_constants.NUM_OUV_BUCKETS
_LABEL_KEY = census_constants.LABEL_KEY
_transformed_name = census_constants.transformed_name

```

```
def _gzip_reader_fn(filenamees):
```

```
    '''Load compressed dataset
```

```
    Args:
```

```
        filenamees - filenamees of TFRecords to load
```

```
    Returns:
```

```
        TFRecordDataset loaded from the filenamees
```

```
    '''
```

```

    # Load the dataset. Specify the compression type since it is saved as `.gz`
    return tf.data.TFRecordDataset(filenamees, compression_type='GZIP')

```

```
def _input_fn(file_pattern,
```

```
              tf_transform_output,
```

```
              num_epochs=None,
```

```
              batch_size=32) -> tf.data.Dataset:
```

```
    '''Create batches of features and labels from TF Records
```

```
    Args:
```

```
        file_pattern - List of files or patterns of file paths containing Example records.
```

```
        tf_transform_output - transform output graph
```

```
        num_epochs - Integer specifying the number of times to read through the dataset.
```

```
            If None, cycles through the dataset forever.
```

```
        batch_size - An int representing the number of records to combine in a single batch.
```

```
    Returns:
```

```
        A dataset of dict elements, (or a tuple of dict elements and label).
```

```
        Each dict maps feature keys to Tensor or SparseTensor objects.
```

```
    '''
```

```

    # Get post-transform feature spec
    transformed_feature_spec = (

```

```
        tf_transform_output.transformed_feature_spec().copy())
```

```

    # Create batches of data
    dataset = tf.data.experimental.make_batched_features_dataset(

```

```
        file_pattern=file_pattern,
```

```
        batch_size=batch_size,
```

```
        features=transformed_feature_spec,
```

```
        reader=_gzip_reader_fn,
```

```
        num_epochs=num_epochs,
```

```
        label_key=_transformed_name(_LABEL_KEY)
```

```
    )
```

```

    return dataset

```

```
def _get_serve_tf_examples_fn(model, tf_transform_output):
```

```
    """Returns a function that receives a serialized TF Example and returns TFEval"""
```

Returns a function that parses a serialized `tf.Example` and applies it.

```
# Get transformation graph
model.tft_layer = tf_transform_output.transform_features_layer()

@tf.function
def serve_tf_examples_fn(serialized_tf_examples):
    """Returns the output to be used in the serving signature."""
    # Get pre-transform feature spec
    feature_spec = tf_transform_output.raw_feature_spec()

    # Pop label since serving inputs do not include the label
    feature_spec.pop(_LABEL_KEY)

    # Parse raw examples into a dictionary of tensors matching the feature spec
    parsed_features = tf.io.parse_example(serialized_tf_examples, feature_spec)

    # Transform the raw examples using the transform graph
    transformed_features = model.tft_layer(parsed_features)

    # Get predictions using the transformed features
    return model(transformed_features)

return serve_tf_examples_fn

def _build_keras_model(hidden_units: List[int] = None) -> tf.keras.Model:
    """Creates a DNN Keras model for classifying income data.

    Args:
        hidden_units: [int], the layer sizes of the DNN (input layer first).

    Returns:
        A keras Model.
    """

    # Use helper function to create the model
    model = _wide_and_deep_classifier(
        dnn_hidden_units=hidden_units or [100, 70, 50, 25])

    return model

def _wide_and_deep_classifier(dnn_hidden_units):
    """Build a simple keras wide and deep model using the Functional API.

    Args:
        wide_columns: Feature columns wrapped in indicator_column for wide (linear)
            part of the model.
        deep_columns: Feature columns for deep part of the model.
        dnn_hidden_units: [int], the layer sizes of the hidden DNN.

    Returns:
        A Wide and Deep Keras model
    """
```



```

# Define input layers for numeric keys
input_numeric = [
    tf.keras.layers.Input(name=_transformed_name(colname), shape=(1,), dtype=tf.float32)
    for colname in _NUMERIC_FEATURE_KEYS
]

# Define input layers for vocab keys
input_categorical = [
    tf.keras.layers.Input(name=_transformed_name(colname), shape=(vocab_size + _NUM_OOV_BUCKETS,),
    for colname, vocab_size in _VOCAB_FEATURE_DICT.items()
]

# Define input layers for bucket key
input_categorical += [
    tf.keras.layers.Input(name=_transformed_name(colname), shape=(num_buckets,), dtype=tf.float32)
    for colname, num_buckets in _BUCKET_FEATURE_DICT.items()
]

# Concatenate numeric inputs
deep = tf.keras.layers.concatenate(input_numeric)

# Create deep dense network for numeric inputs
for numnodes in dnn_hidden_units:
    deep = tf.keras.layers.Dense(numnodes)(deep)

# Concatenate categorical inputs
wide = tf.keras.layers.concatenate(input_categorical)

# Create shallow dense network for categorical inputs
wide = tf.keras.layers.Dense(128, activation='relu')(wide)

# Combine wide and deep networks
combined = tf.keras.layers.concatenate([deep, wide])

# Define output of binary classifier
output = tf.keras.layers.Dense(
    1, activation='sigmoid')(combined)

# Setup combined input
input_layers = input_numeric + input_categorical

# Create the Keras model
model = tf.keras.Model(input_layers, output)

# Define training parameters
model.compile(
    loss='binary_crossentropy',
    optimizer=tf.keras.optimizers.Adam(lr=0.001),
    metrics='binary_accuracy')

# Print model summary
model.summary()

return model

```

```

# TFX Trainer will call this function.
def run_fn(fn_args: FnArgs):
    """Defines and trains the model.

    Args:
        fn_args: Holds args as name/value pairs. Refer here for the complete attributes:
https://www.tensorflow.org/tfx/api\_docs/python/tfx/components/trainer/fn\_args\_utils/FnArgs#attributes
    """

    # Number of nodes in the first layer of the DNN
    first_dnn_layer_size = 100
    num_dnn_layers = 4
    dnn_decay_factor = 0.7

    # Get transform output (i.e. transform graph) wrapper
    tf_transform_output = tft.TFTransformOutput(fn_args.transform_output)

    # Create batches of train and eval sets
    train_dataset = _input_fn(fn_args.train_files[0], tf_transform_output, 10)
    eval_dataset = _input_fn(fn_args.eval_files[0], tf_transform_output, 10)

    # Build the model
    model = _build_keras_model(
        # Construct layers sizes with exponential decay
        hidden_units=[
            max(2, int(first_dnn_layer_size * dnn_decay_factor**i))
            for i in range(num_dnn_layers)
        ])

    # Callback for TensorBoard
    tensorboard_callback = tf.keras.callbacks.TensorBoard(
        log_dir=fn_args.model_run_dir, update_freq='batch')

    # Train the model
    model.fit(
        train_dataset,
        steps_per_epoch=fn_args.train_steps,
        validation_data=eval_dataset,
        validation_steps=fn_args.eval_steps,
        callbacks=[tensorboard_callback])

    # Define default serving signature
    signatures = {
        'serving_default':
            _get_serve_tf_examples_fn(model,
                                       tf_transform_output).get_concrete_function(
                                           tf.TensorSpec(
                                               shape=[None],
                                               dtype=tf.string,
                                               name='examples')),
    }

```

```
# Save model with signature
model.save(fn_args.serving_model_dir, save_format='tf', signatures=signatures)
```

Writing census\_trainer.py

Now, we pass in this model code to the `Trainer` component and run it to train the model.

**Note:** *You can ignore the Exception ignored in: <function CapturableResourceDeleter.\_\_del\_\_> prompt which generates a long Traceback. This is an internal issue with the Tensorflow version used with this version of TFX. This might pop up here and in the Evaluator component.*

```
from tfx.proto import trainer_pb2

trainer = Trainer(
    module_file=os.path.abspath(_census_trainer_module_file),
    examples=transform.outputs['transformed_examples'],
    transform_graph=transform.outputs['transform_graph'],
    schema=schema_gen.outputs['schema'],
    train_args=trainer_pb2.TrainArgs(num_steps=50),
    eval_args=trainer_pb2.EvalArgs(num_steps=50))
context.run(trainer, enable_cache=False)
```

marital-status_xf (InputLayer)	[(None, 12)]	0	
occupation_xf (InputLayer)	[(None, 20)]	0	
race_xf (InputLayer)	[(None, 10)]	0	
relationship_xf (InputLayer)	[(None, 11)]	0	
workclass_xf (InputLayer)	[(None, 14)]	0	
sex_xf (InputLayer)	[(None, 7)]	0	
native-country_xf (InputLayer)	[(None, 47)]	0	
age_xf (InputLayer)	[(None, 4)]	0	
dense_2 (Dense)	(None, 48)	3408	dense_1[0][0]
concatenate_1 (Concatenate)	(None, 146)	0	education_xf[0][0] marital-status_xf[0][0] occupation_xf[0][0] race_xf[0][0] relationship_xf[0][0] workclass_xf[0][0] sex_xf[0][0] native-country_xf[0][0] age_xf[0][0]
dense_3 (Dense)	(None, 34)	1666	dense_2[0][0]
dense_4 (Dense)	(None, 128)	18816	concatenate_1[0][0]
concatenate_2 (Concatenate)	(None, 162)	0	dense_3[0][0] dense_4[0][0]
dense_5 (Dense)	(None, 1)	163	concatenate_2[0][0]

Total params: 31,723  
 Trainable params: 31,723  
 Non-trainable params: 0

50/50 [=====] - 5s 37ms/step - loss: 0.5807 - binary\_accuracy: 0.69

#### ▼ ExecutionResult at 0x7fb3e4313d50

```
.execution_id      6
.component        ▶ Trainer at 0x7fb3e3ed6ad0
.component.inputs  ['examples']      ▶ Channel of type 'Examples' (1 artifact) at
                                0x7fb3e3de5410
                  ['transform_graph'] ▶ Channel of type 'TransformGraph' (1
                                artifact) at 0x7fb3e3de5990
                  ['schema']        ▶ Channel of type 'Schema' (1 artifact) at
                                0x7fb3e3e1cc90
.component.outputs
```

Let's review the outputs of this component. The `model` output points to the model output itself.

```
# Get `model` output of the component
model_artifact_dir = trainer.outputs['model'].get()[0].uri

# List top-level directory
pp.pprint(os.listdir(model_artifact_dir))

# Get model directory
model_dir = os.path.join(model_artifact_dir, 'Format-Serving')

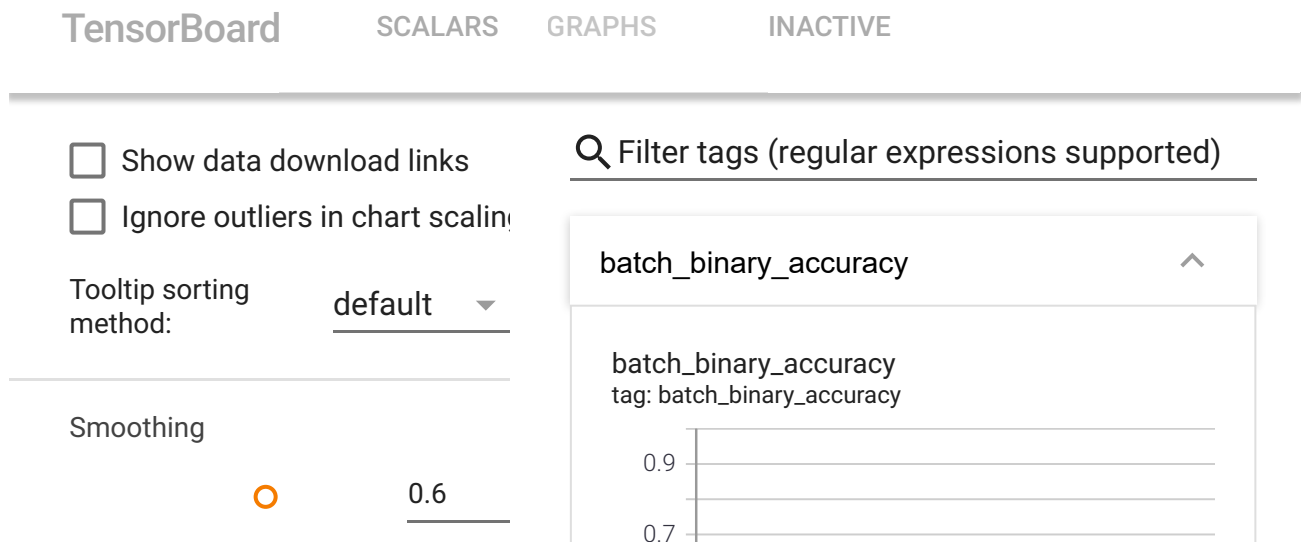
# List subdirectories
pp.pprint(os.listdir(model_dir))

['Format-Serving']
['variables', 'assets', 'saved_model.pb']
```

The `model_run` output acts as the working directory and can be used to output non-model related output (e.g., TensorBoard logs).

```
# Get `model_run` output URI
model_run_artifact_dir = trainer.outputs['model_run'].get()[0].uri

# Load results to Tensorboard
%load_ext tensorboard
%tensorboard --logdir {model_run_artifact_dir}
```



## ▼ Evaluator

The `Evaluator` component computes model performance metrics over the evaluation set using the [TensorFlow Model Analysis](#) library. The `Evaluator` can also optionally validate that a newly trained model is better than the previous model. This is useful in a production pipeline setting where you may automatically train and validate a model every day.

There's a few steps needed to setup this component and you will do it in the next cells.

Write a regex to filter runs

## ▼ Define EvalConfig

First, you will define the `EvalConfig` message as you did in the previous lab. You can also set thresholds so you can compare subsequent models to it. The module below should look familiar. One minor difference is you don't have to define the candidate and baseline model names in the `model_specs`. That is automatically detected.

```
import tensorflow_model_analysis as tfma
from google.protobuf import text_format

eval_config = text_format.Parse("""
## Model information
model_specs {
  # This assumes a serving model with signature 'serving_default'.
  label_key: "label"
}

## Post training metric information
metrics_specs {
  metrics { class_name: "ExampleCount" }
  metrics {
    class_name: "BinaryAccuracy"
    threshold {
      # Ensure that metric is always > 0.5
      value_threshold {
        lower_bound { value: 0.5 }
      }
    }
  }
}
```

```

    }
    # Ensure that metric does not drop by more than a small epsilon
    # e.g. (candidate - baseline) > -1e-10 or candidate > baseline - 1e-10
    change_threshold {
      direction: HIGHER_IS_BETTER
      absolute { value: -1e-10 }
    }
  }
}

metrics { class_name: "BinaryCrossentropy" }
metrics { class_name: "AUC" }
metrics { class_name: "AUCPrecisionRecall" }
metrics { class_name: "Precision" }
metrics { class_name: "Recall" }
metrics { class_name: "MeanLabel" }
metrics { class_name: "MeanPrediction" }
metrics { class_name: "Calibration" }
metrics { class_name: "CalibrationPlot" }
metrics { class_name: "ConfusionMatrixPlot" }
# ... add additional metrics and plots ...
}

## Slicing information
slicing_specs {} # overall slice
slicing_specs {
  feature_keys: ["race"]
}
slicing_specs {
  feature_keys: ["sex"]
}
"""", tfma.EvalConfig())

```

## ▼ Resolve latest blessed model

If you remember in the last lab, you were able to validate a candidate model against a baseline by modifying the `EvalConfig` and `EvalSharedModel` definitions. That is also possible using the `Evaluator` component and you will see how it is done in this section.

First thing to note is that the `Evaluator` marks a model as `BLESSED` if it meets the metrics thresholds you set in the eval config module. You can load the latest blessed model by using the [LatestBlessedModelStrategy](#) with the [Resolver](#) component. This component takes care of finding the latest blessed model for you so you don't have to remember it manually. The syntax is shown below.

```

from tfx.dsl.components.common.resolver import Resolver
from tfx.dsl.input_resolution.strategies.latest_blessed_model_strategy import LatestBlessedModelStrategy
from tfx.types import Channel
from tfx.types.standard_artifacts import Model, ModelBlessing

# Setup the Resolver node to find the latest blessed model
model_resolver = Resolver(
    strategy_class=LatestBlessedModelStrategy,

```

```
model=Channel(type=Model),
model_blessing=Channel(
    type=ModelBlessing)).with_id(
    'latest_blessed_model_resolver')
```

```
# Run the Resolver node
context.run(model_resolver)
```

▼ **ExecutionResult** at 0x7fb3e3eadf10

```
.execution_id      7
.component         <tfx.dsl.components.common.resolver.Resolver object at
                  0x7fb397467a50>
.component.inputs  ['model']          ► Channel of type 'Model' (0 artifacts) at
                                      0x7fb397954050
                  ['model_blessing'] ► Channel of type 'ModelBlessing' (0 artifacts)
                                      at 0x7fb397467090
.component.outputs ['model']          ► Channel of type 'Model' (0 artifacts) at
```

As expected, the search yielded 0 artifacts because you haven't evaluated any models yet. You will run this component again in later parts of this notebook and you will see a different result.

```
# Load Resolver outputs
model_resolver.outputs['model']
```

▼ **Channel** of type 'Model' (0 artifacts) at 0x7fb3e4313850

```
.type_name Model
._artifacts []
```

## ▼ Run the Evaluator component

With the `EvalConfig` defined and code to load the latest blessed model available, you can proceed to run the `Evaluator` component.

You will notice that two models are passed into the component. The `Trainer` output will serve as the candidate model while the output of the `Resolver` will be the baseline model. While you can indeed run the `Evaluator` without comparing two models, it would likely be required in production environments so it's best to include it. Since the `Resolver` doesn't have any results yet, the `Evaluator` will just mark the candidate model as `BLESSED` in this run.

Aside from the eval config and models (i.e. `Trainer` and `Resolver` output), you will also pass in the `raw` examples from `ExampleGen`. By default, the component will look for the `eval` split of these examples and since you've defined the serving signature, these will be transformed internally before feeding to the model inputs.

*Note: You can ignore the Exception ignored in: <function CapturableResourceDeleter.\_\_del\_\_> prompt which generates a long Traceback. This is an internal issue with the Tensorflow version used with this version of TFX. This might pop up here and in the Evaluator component.*



```
# Setup and run the Evaluator component
evaluator = Evaluator(
    examples=example_gen.outputs['examples'],
    model=trainer.outputs['model'],
    baseline_model=model_resolver.outputs['model'],
    eval_config=eval_config)
context.run(evaluator, enable_cache=False)
```

```

File /usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py, line 50:
    g._add_function(self)
File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/framework/ops.py", line 339:
    gradient)
tensorflow.python.framework.errors_impl.InvalidArgumentError: 'func' argument to TF_GraphCopy
Exception ignored in: <function CapturableResourceDeleter.__del__ at 0x7fb408066050>
Traceback (most recent call last):
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/training/tracking/tracking.py", line 100:
    self._destroy_resource()
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py", line 100:
    result = self._call(*args, **kwargs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py", line 100:
    self._initialize(args, kwargs, add_initializers_to=initializers)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py", line 100:
    *args, **kwargs))
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 29:
    graph_function, _ = self._maybe_define_function(args, kwargs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 33:
    graph_function = self._create_graph_function(args, kwargs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 32:
    capture_by_value=self._capture_by_value),
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/framework/func_graph.py", line 100:
    func_outputs = python_func(*func_args, **func_kwargs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/def_function.py", line 100:
    out = weak_wrapped_fn().__wrapped__(*args, **kwargs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/saved_model/function_deserializer.py", line 100:
    return _call_concrete_function(function, inputs)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/saved_model/function_deserializer.py", line 100:
    result = function._call_flat(tensor_inputs, function._captured_inputs) # pylint: disable=protected-access
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/saved_model/load.py", line 100:
    cancellation_manager)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 19:
    flat_outputs = forward_function.call(ctx, args_with_tangents)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 58:
    executor_type=executor_type)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/ops/functional_ops.py", line 100:
    f.add_to_graph(graph)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/eager/function.py", line 50:
    g._add_function(self)
  File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/framework/ops.py", line 339:
    gradient)
tensorflow.python.framework.errors_impl.InvalidArgumentError: 'func' argument to TF_GraphCopy
WARNING:root:Make sure that locally built Python SDK docker image has Python 3.7 interpreter

```

#### ▼ ExecutionResult at 0x7fb397371290

<b>.execution_id</b>	8
<b>.component</b>	► <b>Evaluator</b> at 0x7fb395d0c450
<b>.component.inputs</b>	<div> <div>['examples']</div> <div>► Channel of type 'Examples' (1 artifact) at 0x7fb3e561bf50</div> </div> <div> <div>['model']</div> <div>► Channel of type 'Model' (1 artifact) at 0x7fb3e3ed6850</div> </div> <div> <div>['baseline_model']</div> <div>► Channel of type 'Model' (0 artifacts) at 0x7fb3e4313850</div> </div>
<b>.component.outputs</b>	<div> <div>['evaluation']</div> <div>► Channel of type 'ModelEvaluation' (1 artifact) at 0x7fb3e4313850</div> </div>









